

DD2476: Search Engines and Information Retrieval Systems

Johan Boye*

KTH

Lecture 2

* Many slides inspired by Manning, Raghavan and Schütze

Indexing pipeline

Documents



Byte stream

F r i e n d s , r o m a n s , a n d c o u n t r y m e n


Token stream

Friends romans countrymen

Term stream

friend roman countryman

Inverted index



Basic text processing

- Text comes in many **different formats** (html, text, Word, Excel, PDF, PostScript, ...), **languages** and **character sets**
- It might need to be
 - separated from images and other non-textual content
 - stripped of markup in HTML or XML

Character formats

- Text encodings
 - **ASCII** (de-facto standard from 1968), 7-bit (=128 chars, 94 printable). Most common on the www until Dec 2007.
 - **Latin-1** (ISO-8859-1), 8-bit, ASCII + 128 extra chars
 - **Unicode** (109 000 code points)
 - **UTF-8** (variable-length encoding of Unicode)

Tokenization

How many tokens are there in this text?

- Look, harry@hp.com, that's Harry's mail address at Hewlett-Packard. Boy, does that guy know Microsoft Word! He's really working with the state-of-the-art in computers. And yesterday he told me my IP number is 131.67.238.92. :-)

Tokenization

- A token is a **meaningful minimal unit** of text.
- Usually, **spaces and punctuation** delimit tokens
- Is that always the case?
 - **San Francisco, Richard III, et cetera, ...**
 - **J.P. Morgan & co**
 - **<http://www.kth.se>, jboye@nada.kth.se**
 - **:-)**
- The exact definition is application-dependent:
 - Sometimes it's important to include punctuation among the tokens (e.g. language modeling)
 - Sometimes it's better not to (e.g. search engines)

Some tricky tokenization issues

- Apostrophes
 - Finland's → Finland's? Finlands? Finland? Finland s?
 - don't → don't ? don t ? do not ? don t?
- Hyphens
 - state-of-the-art → state-of-the-art? state of the art?
 - Hewlett-Packard
 - the *San Francisco-Los Angeles* flight
- Numbers
 - Can contain spaces or punctuation: **123 456.7** or **123,456.7** or **123 456,7**
 - **+46 (8) 790 60 00**
 - **131.169.25.10**
 - My PGP key is **324a3df234cb23e**

So how do we do it?

- In assignment 1.1:
 - In the general case, assume that space and punctuation (except apostrophes and hyphens) separate tokens
 - **Specify special cases with regular expressions**

Normalization

- After tokenization, we sometimes need to “normalize” tokens
 - Abbreviations: **U.S., US** → **U.S.**
 - Case folding: **Window, window** → **window**
 - Diacritica: **a, å, ä, à, á, â** → **a, c, ç, č** → **c, n, ñ** → **n, l, t, → l, ...**
 - Umlaut: **Tübingen** → **Tuebingen**, **Österreich** → **Oesterriegch**
- Need for normalization is highly dependent on application
 - Is it always a good idea to lowercase Apple and Windows?
 - Should we remove diacritica?
 - When should we regard run and runs as the same word?

Morphemes


- Words are built from smaller meaningful units called **morphemes**.
- A morpheme belongs to one of two classes:
 - **stem**: the core meaning-bearing unit
 - **affix**: small units glued to the stem to signal various grammatical functions
- An affix can in its turn be classified as a
 - **prefix** (*un-*)
 - **suffix** (*-s, -ed, -ly*)
 - **infix** (*Swedish korru-m-pera*)
 - **circumfix** (*German ge-sag-t*)

Word formation

- Words can be **inflected** to signal grammatical information:
 - play, plays, played, playing
 - cat, cats, cat's, cats'
- Words can also be **derived** from other words:
 - friend → friendly → friendliness → unfriendliness
- Words can be **compound**:
 - smart + phone → smartphone
 - anti + missile → anti-missile
- Clitics
 - Le + hôtel → L'hôtel, Ce + est → c'est
 - She is → she's, She has → she's

Language variation

- English morphology is exceptionally simple!



Language variation

Parler

The verb *parler* "to speak", in French orthography and IPA transcription

	Indicative					Subjunctive		Conditional	Imperative
	Present	Simple past	Imperfect	Simple future	Present	Imperfect	Present	Present	Present
je	parl-e /paʁl/	parl-al /paʁla/	parl-ais /paʁlɛ/	parl-eral /paʁləʁ/	parl-e /paʁl/	parl-assee /paʁlaʁ/	parl-erais /paʁləʁ/		
tu	parl-es /paʁl/	parl-as /paʁla/	parl-ais /paʁlɛ/	parl-eras /paʁləʁa/	parl-es /paʁl/	parl-asses /paʁlaʁ/	parl-erais /paʁləʁ/	parl-e /paʁl/	
Il	parl-e /paʁl/	parl-a /paʁla/	parl-alt /paʁlɛ/	parl-era /paʁləʁa/	parl-e /paʁl/	parl-ât /paʁla/	parl-erait /paʁləʁa/		
nous	parl-ons /paʁlɔ̃/	parl-âmes /paʁlɑ̃m/	parl-ions /paʁlɔ̃/	parl-erons /paʁləʁɔ̃/	parl-ions /paʁlɔ̃/	parl-assions /paʁlaʃɔ̃/	parl-erions /paʁləʁɔ̃/	parl-ons /paʁlɔ̃/	
vous	parl-ez /paʁle/	parl-âtes /paʁlɑ̃t/	parl-ieez /paʁlje/	parl-erez /paʁləʁe/	parl-lez /paʁlje/	parl-assiez /paʁlaʃje/	parl-eriez /paʁləʁje/	parl-ez /paʁle/	
Ils	parl-ent /paʁl/	parl-èrent /paʁləʁ/	parl-aient /paʁlɛ/	parl-eront /paʁləʁɑ̃/	parl-ent /paʁl/	parl-assent /paʁlaʃ/	parl-eraien /paʁləʁɑ̃/		

Some non-English words

- German: ***Lebensversicherungsgesellschaftsangestellter***
 - “Life insurance company employee”
- Greenlandic: ***iglukpisuktunga***
 - *iglu* = house, *kpi* = build, *suk* = (I) want, *tu* = myself, *nga* = me
- Finnish: ***järjestelmättömyydellänsäkäänköhän***
 - “not even with its lack of order”

Lemmatization

- Map **inflected form** to its **lemma** (=base form)
- “The boys’ cars are different colours” → “The boy car be different color”
- Requires language-specific linguistic analysis
 - part-of-speech tagging
 - morphological analysis
- Particularly useful in morphologically rich languages, like Finnish, Turkish, Hungarian

Stemming


- Don't do morphological or syntactic analysis, just **chop off the suffixes**
 - No need to know that "foxes" is plural of "fox"
- Much **less expensive** than lemmatization, but **can be very wrong** sometimes
 - stocks → stock, stockings → stock
- Stemming usually improves **recall** but lowers **precision**

Porter's algorithm

- Rule-based stemming for English
 - ATIONAL → ATE
 - SSES → SS
 - ING → ε
- Some context-sensitivity
- (W>1) EMENT → ε
 - REPLACEMENT → REPLAC
 - CEMENT → CEMENT

Compound splitting

Can be achieved with finite-state techniques.



Compound splitting

- In Swedish: **försäkringsbolag** (insurance company)
 - **bolag** is the head
 - **försäkring** is a modifier
 - the **s** is an infix
- This process can be recursive:
 - försäkringsbolagslagen (the insurance company law)
 - **en** is a suffix indicating definite form
 - **lag** is the head
 - the **s** is an infix
 - **försäkringsbolag** is the modifier

Stop words

- Can we exclude the most common words?
 - In English: **the, a, and, to, for, be, ...**
 - Little semantic content
 - ~30% of postings for top 30 words
- However:
 - "**Let it be**", "**To be or not to be**", "**The Who**"
 - "**King of Denmark**"
 - "**Flights to London**" vs "**Flights from London**"
 - Trend is to keep stop words: compression techniques means that space requirements are small

Language-specific issues

- Chinese and Japanese have no spaces between words:
 - 莎拉波娃现在居住在美国东南部的佛罗里达。
 - Not always guaranteed a unique tokenization
- Japanese have several alphabets
 - **Katakana** and **Hiragana** (syllabic)
 - **Kanji** (Chinese characters)
 - **Romaji** (Western characters)
 - All of these may be intermingled in the same sentence

Chinese tokenization

我喜欢新西兰花

Chinese tokenization

我喜欢新西兰花

我 | 喜欢 | 新西兰 | 花
“I like New Zealand flowers”

我 | 喜欢 | 新 | 西兰花
“I like fresh broccoli”

Chinese tokenization

The greedy matching algorithm:

1. Put a pointer in the beginning of the string
2. Find the longest prefix of the string that matches a word in the dictionary
3. Move the pointer over that prefix
4. Go to 2

我喜欢新西兰花

我 | 喜欢 | 新西兰 | 花
“I like New Zealand flowers”

我 | 喜欢 | 新 | 西兰花
“I like fresh broccoli”

Greedy matching

The cat in the hat → The cat in the hat

The table down there → ?

- Wouldn't work so well for English
- But works very well for Chinese

Sum-up


- **Reading, tokenizing and normalizing** contents of documents
 - **File types and character encodings**
 - Tokenization issues: **punctuation, compound words, word order, stop words**
 - Normalization issues: **diacritica, case folding, lemmatization, stemming**
- We're ready for **indexing**

Indexing and search

- **Recap:**
 - We want to quickly find the **most relevant documents** satisfying our **information need**.
 - The user gives a **search query**.
 - The engine searches through the **index**, retrieves the **matching** documents, and possibly **ranks** them.

The index

- Conceptually: the **term-document matrix**



One-word queries

denmark

- Return all the documents in which '**denmark**' appears. (Task 1.2)

Multi-word queries

copenhagen denmark


- **Intersection** query (Task 1.3)
- **Phrase** query (Task 1.4)
- **Union** query (Assignment 2)

Practical indexing

- We need a **sparse matrix representation**.
- In the computer assignments we use:
 - a **hashtable** for the dictionary
 - **sorted arraylists** for the rows
- Rows are called **postings lists**.

Intersection


- Walk through two postings lists simultaneously



- Runs in $O(n+m)$, where n,m are the lengths of the lists
- We can do better (if index isn't changing too fast)

Skip pointers


- Add skip pointers at indexing time



- By using skip pointers, we don't have to compare 41 to 17 or 21

Skip pointers: Where?

- Tradeoff:
 - More skips → shorter skip spans \Rightarrow more likely to skip.
But lots of comparisons to skip pointers.
 - Fewer skips → few pointer comparison, but then long skip spans \Rightarrow few successful skips.
 - Heuristic: for length L , use \sqrt{L} evenly spaced skip pointers



Phrase queries

- E.g. **“Donald Trump”**
- Should not match “President Trump”
 - The concept of phrase queries has proven easily understood by users; one of the few “advanced search” ideas that works
 - Many more queries are *implicit phrase queries*
- For this, it no longer suffices to store only $\langle \text{term} : \text{docs} \rangle$ entries

First attempt: Biword index

- “Friends, Romans, Countrymen” generates the **biwords**
 - *friends romans*
 - *romans countrymen*
- Each of these biwords is now a dictionary term
- Two-word phrase query-processing is now immediate.
- Longer phrases: **friends romans countrymen**
- Intersect **friends romans** and **romans countrymen?**

Biword index: disadvantages

- **False positives**
 - Requires post-processing to avoid
- **Index blowup** due to bigger dictionary
 - Infeasible for more than biwords, big even for them

Positional indexes

- For each term and doc, store the positions where (tokens of) the term appears

<be;

1: 7, 18, 33, 72, 86, 231;

2: 3, 149;

4: 17, 191, 291, 430, 434;

5: 363, 367, ...>

- Intersection needs to deal with more than equality

Processing phrase queries

- Extract inverted index entries for each distinct term: ***to, be, or, not.***
- Intersect their *doc:position* lists to enumerate all positions with “***to be or not to be***”.
 - ***to:***
 - 2:1,17,74,222,551; **4:8,16,190,429,433;** 7:13,23,191; ...
 - ***be:***
 - 1:17,19; **4:17,191,291,430,434;** 5:14,19,101; ...
- Same general method for proximity searches

Exercise

Which docs match the query "**fools rush in**" ?

fools: 2: 1,17,74,222;
4: 78,108,458;
7: 3,13,23,193;

in: 2: 3,37,76,444,851;
4: 10,20,110,470,500;
7: 5,15,25,195;

rush: 2: 2,75,194,321,702;
4: 9,69,149,429,569;
7: 14,404;


Positional index size

- Need an entry for each occurrence, not just once per document
- Consider a term with frequency 0.1%
 - Doc contain 1000 tokens → 1 occurrence
 - 100 000 tokens → 100 occurrences
- Rule of thumb: is 2–4 as large as a non-positional index
- Positional index size 35–50% of volume of original text
- Caveat: all of this holds for “English-like” languages


Dictionary structures

- How do we store terms in the dictionary?
- Hash tables:
 - **Lookup in constant time** $O(1)$
 - **No wildcard queries**
 - Occasionally we need to **rehash everything** as the vocabulary grows. This is **expensive**.
- Trees:
 - **Lookup in logarithmic time** (if tree is **balanced**)
 - Allows for **wildcard queries**
 - Requires standard (alphabetical) **order of terms**

Binary tree




B-tree



Every internal node has a number of children in the interval $[a,b]$ where a, b are appropriate natural numbers, e.g., $[2,4]$.


Large indexes (Task 1.7-1.8)

- The web is big:
 - 1998: **26 million** unique web pages
 - 2018: **130 trillion** (1.3×10^{14}) unique web pages!
 - about 4.26 billion of these are indexed.
- What if the index is **too large to fit in main memory?**
 - Dictionary in main memory, postings on disk
 - Dictionary (partially) on disk, postings on disk
 - Index on several disks (cluster)



Mixed MM and disk storage

- **Tree:** Nodes below a certain depth stored on disk



- **Hashtable:** All postings put on disk, hash keys in MM
- A **distributed hash table** allows keys and postings to be distributed over a large number of computers

Hardware basics

- Access to data in memory is ***much*** faster than access to data on disk.
- **Disk seeks:** No data is transferred from disk while the disk head is being positioned.
 - Therefore: Transferring **one large chunk of data** from disk to memory is faster than transferring many small chunks.
 - Disk I/O is **block-based**: Reading and writing of entire blocks (as opposed to smaller chunks).
 - Block sizes: 8KB to 256 KB.

Hardware assumptions

- In the book:

statistic

average seek time

value
$$5 \text{ ms} = 5 \times 10^{-3} \text{ s}$$

transfer time per byte

$$0.02 \mu\text{s} = 2 \times 10^{-8} \text{ s}$$

processor's clock rate

$$10^9 \text{ s}^{-1} \text{ (1 GHz)}$$

low-level operation

$$0.01 \mu\text{s} = 10^{-8} \text{ s}$$

(e.g., compare & swap a word)

size of main memory

several GB

size of disk space

1 TB or more

Basic indexing

- Term-document pairs are collected when documents are parsed

Doc 1

I did enact Julius
Caesar I was killed
i' the Capitol;
Brutus killed me.

Doc 2

So let it be with
Caesar. The noble
Brutus hath told you
Caesar was ambitious



Term	Doc #
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2

Sorting step

- The list of term-doc pairs is sorted
- This must be done on disk for large lists
- Goal: Minimize the number of disk seeks


Term	Doc #	Term	Doc #
I	1	ambitious	2
did	1	be	2
enact	1	brutus	1
julius	1	brutus	2
caesar	1	capitol	1
I	1	caesar	1
was	1	caesar	2
killed	1	caesar	2
I'	1	did	1
the	1	enact	1
capitol	1	hath	1
brutus	1	I	1
killed	1	I	1
me	1	I'	1
so	2	it	2
let	2	julius	1
it	2	→ killed	1
be	2	killed	1
with	2	let	2
caesar	2	me	1
the	2	noble	2
noble	2	so	2
brutus	2	the	1
hath	2	the	2
told	2	told	2
you	2	you	2
caesar	2	was	1
was	2	was	2
ambitious	2	with	2

A bottleneck

- Say we want to sort 100,000,000 term-doc-pairs.
- A list can be sorted by $N \log_2 N$ comparison operations.
 - How much time does that take? (assume 10^{-8} s/operation)
- Suppose that each comparison additionally took 2 disk seeks
 - How much time? (assuming 5×10^{-3} /disk seek)

Blocked sort-based indexing

- (term-doc) records
 - Define a Block ~ 10M of such records
 - Accumulate postings for each block, sort, write to disk.
 - Then merge the blocks into one long sorted order.



Sorting 10 blocks of 10M records

- First, read each block and sort within:
 - Quicksort takes $N \log_2 N$ expected steps
 - In our case (10M) $\log_2 (10M)$ steps
- *Exercise: estimate total time to read each block from disk and quicksort it.*
 - assuming transfer time 2×10^{-8} s per byte
- 10 times this estimate – gives us 10 sorted runs of 10M records each.

Blocked sort-based indexing

BSBINDEXCONSTRUCTION()

- 1 $n \leftarrow 0$
- 2 **while** (all documents have not been processed)
- 3 **do** $n \leftarrow n + 1$
- 4 $block \leftarrow \text{PARSENEXTBLOCK}()$
- 5 BSBI-INVERT($block$)
- 6 WRITEBLOCKToDISK($block, f_n$)
- 7 MERGEBLOCKS($f_1, \dots, f_n; f_{\text{merged}}$)

From BSBI to SPIMI

- BSBI requires that the dictionary can be kept in main memory
- Alternative approach: Construct several **separate** indexes and **merge** them
 - Generate separate dictionaries for each block
 - No need to keep dictionary in main memory
 - Accumulate postings directly in postings list (as in assignment 1).
- This is called **SPIMI** – Single-Pass In-Memory Index construction (Figure 4.4)

SPIMI-invert


```
SPIMI-INVERT(token_stream)
```

```
1  output_file = NEWFILE()
2  dictionary = NEWHASH()
3  while (free memory available)
4  do token  $\leftarrow$  next(token_stream)
5    if term(token)  $\notin$  dictionary
6      then postings_list = ADDToDICTIONARY(dictionary, term(token))
7      else postings_list = GETPOSTINGSLIST(dictionary, term(token))
8      if full(postings_list)
9        then postings_list = DOUBLEPOSTINGSLIST(dictionary, term(token))
10       ADDTOPSTINGSLIST(postings_list, docID(token))
11  sorted_terms  $\leftarrow$  SORTTERMS(dictionary)
12  WRITEBLOCKTODISK(sorted_terms, dictionary, output_file)
13  return output_file
```

Large indexes (Task 1.7-1.8)


- Task 1.7 asks you to implement an index which is stored on disk
 - using any method (well, not quite...) for grade C
 - using a hash table with both dictionary and postings lists on disk for grade B
- Task 1.8 ask you to to a kind of SPIMI-invert, for grade A
 - construction of partial hash tables
 - merging of the partial hash tables into the final table

Hash tables on disk- what one would like to do




Why doesn't this work?

Hash tables on disk- what we will do





Hash table on disk

- Dictionary file:
 - with entries of a **fixed length**
 - entries contain pointer to the data file
- Data file
 - contains string representation of postings list
 - don't serialize the PostingsList objects! (waste of space)
- Hash function
 - inputs word (as a string)
 - outputs an integer [0...TABLESIZE-1] which is a pointer to the dictionary file.



"words are useful"







Hash table on disk

- Dictionary file
 - has a fixed size
 - will be mostly empty (load factor about 0.33)
- Data file grows dynamically
 - will be completely packed


Hash collisions



Hash collisions



Hash collisions



Hash function

- Have a look in the literature
 - or devise your own method
 - but be sure there aren't too many collisions
 - about 1 collision/unique word is a reasonable target
 - (that means about 200,000 collisions)

Dynamic indexing

- Up to now, we have assumed that collections are static.
- They rarely are:
 - Documents come in over time and need to be inserted.
 - Documents are deleted and modified.
- This means that the dictionary and postings lists have to be modified:
 - Postings updates for terms already in dictionary
 - New terms added to dictionary

Simplest approach

- Maintain “big” main index
- New docs go into “small” auxiliary index
- Search across both, merge results
- Deletions
 - Invalidation bit-vector for deleted docs
 - Filter docs output on a search result by this invalidation bit-vector
- Periodically, re-index into one main index

Logarithmic merge

- Maintain a series of indexes, each twice as large as the previous one.
- Keep smallest (Z_0) in memory
- Larger ones (I_0, I_1, \dots) on disk
- If Z_0 gets too big ($> n$), write to disk as I_0
 - or merge with I_0 (if I_0 already exists) as Z_1
- Either write merge Z_1 to disk as I_1 (if no I_1)
 - or merge with I_1 to form Z_2
- etc.

Dynamic indexing at search engines

- All the large search engines now do dynamic indexing
- Their indices have frequent incremental changes
 - News items, blogs, new topical web pages
- But (sometimes/typically) they also periodically reconstruct the index from scratch
 - Query processing is then switched to the new index, and the old index is then deleted

Assignment 1

- Tokenization (1.1)
- Basic indexing (1.2)
- Intersection search (1.3)
- Phrase search (1.4)
- Evaluation (1.5)
- Query construction (1.6)
- Large indexes on disk (1.7)
- Merging indexes (1.8)