

DD2476: Search Engines and Information Retrieval Systems

Johan Boye*

KTH


Lecture 5

* Many slides inspired by Manning, Raghavan and Schütze

Static quality scores

- We want top-ranking documents to be both **relevant** and **authoritative**
 - Relevance – cosine scores
 - Authority – query-independent property
- Assign **query-independent quality score** $g(d)$ in $[0,1]$ to each document d
- $\text{net-score}(q,d) = w_1 \times g(d) + w_2 \times \cos(q,d)$

The Web as a directed graph



Using link structure for ranking

- **Assumption:** A link from X to Y signals that X's author perceives Y to be an authoritative page.
 - X “casts a vote” on Y.
- **Simple suggestion:** Rank = number of in-links
- However, there are problems with this naive approach.

PageRank: basic ideas

- WWW's particular structure can be exploited
 - pages have links to one another
 - the more in-links, the higher rank
 - in-links from pages having **high rank** are **worth more** than links from pages having low rank
 - this idea is the cornerstone of **PageRank** (Brin & Page 1998)

The random surfer model

- Imagine a **random surfer** that follow links
- The link to follow is selected with uniform probability
- If the surfer reaches a **sink** (a page without links), he **randomly restarts** on a new page
- Every once in a while, the surfer **jumps to a random page** (even if there are links to follow)



PageRank


- With **probability $1-c$** the surfer is bored, **stops following links**, and **restarts** on a random page
- Guess: Google uses $c=0.85$

$$PR(p) = c \left(\sum_{q \in in(p)} \frac{PR(q)}{L_q} \right) + \frac{(1 - c)}{N}$$

- Without this assumption, the surfer will get stuck in a subset of the web.

PageRank example

$$PR_4 = 0.85 \cdot \left(\frac{PR_2}{2} + PR_3 \right) + \frac{0.15}{5}$$



$$PR_3 = 0.85 \cdot \left(\frac{PR_0}{3} + PR_1 + \frac{PR_2}{2} + \frac{PR_4}{5} \right) + \frac{0.15}{5}$$

$$PR_2 = PR_1 = 0.85 \cdot \left(\frac{PR_0}{3} + \frac{PR_4}{5} \right) + \frac{0.15}{5}$$


$$PR_0 = 0.85 \cdot \left(\frac{PR_4}{5} \right) + \frac{0.15}{5}$$

PageRank- interpretations

- Authority / popularity / relative information value
- PR_p = the probability that the random surfer will be at page p at any given point in time
- This is called the **stationary probability**
- How do we compute it?

Random surfer as a Markov chain


- The random surfer model suggests a a **Markov chain** formulation
 - A Markov chain consists of n states, plus an $n \times n$ transition probability matrix P .
 - At each step, we are in exactly one of the states.
 - For $1 \leq i, j \leq n$, the matrix entry P_{ij} tells us the probability of j being the next state, given we are currently in state i .



Transition matrices

P

0	0.33	0.33	0.33	0
0	0	0	1	0
0	0	0	0.5	0.5
0	0	0	0	1
0.2	0.2	0.2	0.2	0.2



$$G = cP + (1-c)J$$

J

0.2	0.2	0.2	0.2	0.2
0.2	0.2	0.2	0.2	0.2
0.2	0.2	0.2	0.2	0.2
0.2	0.2	0.2	0.2	0.2
0.2	0.2	0.2	0.2	0.2

0.0300	0.3105	0.3105	0.3105	0.0300
0.0300	0.0300	0.0300	0.8800	0.0300
0.0300	0.0300	0.0300	0.4550	0.4550
0.0300	0.0300	0.0300	0.0300	0.8800
0.2000	0.2000	0.2000	0.2000	0.2000

Ergodic Markov chains


- A Markov chain is **ergodic** if
 - you have a path from any state to any other
 - For any start state, after a finite transient time T_0 , **the probability of being in any state at a fixed time $T > T_0$ is nonzero.**
- Our transition matrix is non-zero positive everywhere
 \leftrightarrow
the graph is strongly connected \leftrightarrow
the Markov chain is ergodic \leftrightarrow
unique stationary probabilities exist

Probability vectors

- A **probability (row) vector** $\mathbf{x} = (x_1, \dots x_n)$ tells us where the walk is at any point.
- E.g., $(000\dots 1\dots 000)$ means we're in state i .

1 i n

- More generally, the vector $\mathbf{x} = (x_1, \dots x_n)$ means the walk is in state i with probability x_i .
- \mathbf{xG} gives the next time step.
- \mathbf{x} is stationary if $\mathbf{x}=\mathbf{xG}$

$$\sum_{i=1}^n x_i = 1.$$


Stationary probabilities

- Let $\mathbf{a} = (a_1, \dots, a_n)$ denote the row vector of stationary probabilities.
 - If our current position is described by \mathbf{a} , then the next step is distributed as $\mathbf{a}\mathbf{G}$.
 - But \mathbf{a} is the steady state, so $\mathbf{a}=\mathbf{a}\mathbf{G}$.
- Solving the matrix equation $\mathbf{x}=\mathbf{x}\mathbf{G}$ gives us \mathbf{a} .
 - So \mathbf{a} is the (left) eigenvector for \mathbf{G} .
- \mathbf{x} can be computed using **power iteration**.

Power iteration

- Start with any distribution (say $\mathbf{x}=(10\dots 0)$).
 - After one step, we're at \mathbf{xG} ;
 - after two steps at $(\mathbf{xG})\mathbf{G}$, then $((\mathbf{xG})\mathbf{G})\mathbf{G}$ and so on.
- “Eventually”, for “large” k , $\mathbf{xG}^k = \mathbf{a}$.


Power iteration algorithm

```
Let x=(0,...,0) and x' an initial state, say (1,0,...,0)
```

```
while ( |x-x'| > ε ) :  
    x = x'  
    x' = xG
```


This algorithm converges **very** slowly.

Example




	t=0	t=1
0	1	.030
1	0	.313
2	0	.313
3	0	.313
4	0	.030

Example




	t=0	t=1	t=2
0	1	.030	.035
1	0	.313	.044
2	0	.313	.044
3	0	.313	.443
4	0	.030	.435

Example



	t=0	t=1	t=2	t=3	t=4	t=5	...	t=16
0	1	.030	.035	.104	.115	.082095
1	0	.313	.044	.114	.144	.115122
2	0	.313	.044	.114	.144	.115122
3	0	.313	.443	.169	.289	.299278
4	0	.030	.435	.499	.307	.390383

Example



- Stationary probabilities:
(0.095, 0.122, 0.122, 0.278, 0.383)

found after 16 iterations starting from
(1, 0, 0, 0, 0)

Example

- Stationary probabilities:

$$x = (0.095, 0.122, 0.122, 0.278, 0.383)$$

- Check: $xG =$

$$(0.095 \quad 0.122 \quad 0.122 \quad 0.278 \quad 0.383)$$

$$\begin{bmatrix} 0.0300 & 0.3105 & 0.3105 & 0.3105 & 0.0300 \\ 0.0300 & 0.0300 & 0.0300 & 0.8800 & 0.0300 \\ 0.0300 & 0.0300 & 0.0300 & 0.4550 & 0.4550 \\ 0.0300 & 0.0300 & 0.0300 & 0.0300 & 0.8800 \\ 0.2000 & 0.2000 & 0.2000 & 0.2000 & 0.2000 \end{bmatrix}$$

$$= (0.0951 \quad 0.1218 \quad 0.1218 \quad 0.2773 \quad 0.3833)$$

- So x is (approximately) the left eigenvector of G .

Approximating PageRank

- For large graphs, power iteration takes a loooong time.
- We can also approximate PageRank by **simulating the random surfer**.

Monte Carlo simulation to approximate PageRank

- Avrachenkov et al. describe a number of simulation algorithm based on the random surfer model
- D = document id
- Consider a random walk $\{D_t\}_{t \geq 0}$ that starts from some page.
- At each step t :
 - Prob c : D_t = one of the documents with edges from D_{t-1}
 - Prob $(1 - c)$: The random walk terminates

1. MC end-point with random start

- Simulate N runs of the random walk $\{D_t\}_{t \geq 0}$ initiated at a **randomly chosen page**
- PageRank of page $j = 1, \dots, n$:
$$\pi_j = (\#\text{walks which end at } j)/N$$
- **Worst case:** $N = O(n^2)$
- **Mean case:** $N = O(n)$

2. MC end-point with cyclic start

- Simulate $N = mn$ runs of the random walk $\{D_t\}_{t \geq 0}$ initiated at **each page exactly m times**
- PageRank of page $j = 1, \dots, n$:
$$\pi_j = (\text{\#walks which end at } j)/N$$

4. MC complete path stopping at dangling nodes

- Simulate $N = mn$ runs of the random walk $\{D_t\}_{t \geq 0}$ initiated at **each page exactly m times** and **stopping when it reaches a dangling node**
- PageRank of page $j = 1, \dots, n$:
$$\pi_j = (\text{\#visits to node } j \text{ during walks}) / (\text{total \#visits during walks})$$

5. MC complete path with random start

- Simulate N runs of the random walk $\{D_t\}_{t \geq 0}$ initiated at a **randomly chosen page** and **stopping when it reaches a dangling node**
- PageRank of page $j = 1, \dots, n$:
$$\pi_j = (\text{\#visits to node } j \text{ during walks}) / (\text{total \#visits during walks})$$

Monte Carlo advantages

- **Power iteration:**
 - precision improves linearly for all docs
 - computationally expensive
 - must be redone when new pages are added
- **Monte Carlo method:**
 - precision improves faster for high-rank docs
 - parallel implementation possible
 - continuous update

Assignment 2

- 2.5 Implement power iteration
Compute pageranks of links1000.txt
- 2.6 Power iteration with sparse matrices
Compute pageranks of linksDavis.txt
Combine tf-idf and pagerank
- 2.7 Implement Monte-Carlo pagerank computations
Compute pageranks of linksWiki.txt
- 2.8 Hubs and authorities

HITS (Hypertext-Induced Topic Selection)

- HITS is an method similar to PageRank
 - proposed the same year (1998)
- Goal is to find high-quality pages of two kinds:
 - **Hubs**: linking to pages with great authority
 - **Authorities**: linked from great hubs


A typical hub

culture trip

Explore your world

Sections Wis

Home > Sweden > Stockholm >



Stockholm / BARS & CAFES



The 10 Best Cafes And Coffeehouses In Stockholm


Hubs and authorities

- Each page i has two scores:
 - authority score a_i
 - hub score h_i
- HITS algorithm: Initialize $a_i = 1$ $h_i = 1$
- Then iterate until convergence:

– Authority: $\forall i : a_i := \sum_{k \rightarrow i} h_k$

– Hub: $\forall i : h_i := \sum_{i \rightarrow k} a_k$

– Normalize $\sqrt{\sum_i a_i^2} := 1$ $\sqrt{\sum_i h_i^2} := 1$



Matrix formulation


- $a = (a_1, \dots, a_n)$ $h = (h_1, \dots, h_n)$ are authority/hub scores for n documents
- A is the **adjacency matrix**: $A_{ik} = 1$ iff $i \rightarrow k$
- Then do:

$$h_i := \sum_{i \rightarrow k} a_k \Leftrightarrow h_i := \sum_k A_{ik} a_k$$

- i.e. $h := Aa$
- and likewise $a := A^T h$

HITS algorithm (example)

$$A = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{bmatrix}$$




a(1)	1					
a(2)	1					
a(3)	1					

h(1)	1					
h(2)	1					
h(3)	1					

HITS algorithm (example)

$$A = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{bmatrix}$$



$$a = A^T h$$


a(1)	1	2			
a(2)	1	2			
a(3)	1	2			

$$h = A a$$

h(1)	1	3			
h(2)	1	2			
h(3)	1	1			

HITS algorithm (example)

$$A = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{bmatrix}$$



a(1)	1	2	$2/\sqrt{12}=0.577$		
a(2)	1	2	$2/\sqrt{12}=0.577$		
a(3)	1	2	$2/\sqrt{12}=0.577$		


Normalize

h(1)	1	3	$3/\sqrt{14} = 0.802$		
h(2)	1	2	$2/\sqrt{14}=0.535$		
h(3)	1	1	$1/\sqrt{14}=0.267$		

Normalize

HITS algorithm (example)

$$A = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{bmatrix}$$




a(1)	1	0.577	0.615	a = $A^T h$, then normalize
a(2)	1	0.577	0.492	
a(3)	1	0.577	0.615	

h(1)	1	0.802	0.793	h = $A a$, then normalize
h(2)	1	0.535	0.566	
h(3)	1	0.267	0.226	

HITS algorithm (example)

$$A = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{bmatrix}$$



a(1)	1	0.577	0.615	0.627
a(2)	1	0.577	0.492	0.460
a(3)	1	0.577	0.615	0.627

Convergence

h(1)	1	0.802	0.793	0.789
h(2)	1	0.535	0.566	0.577
h(3)	1	0.267	0.226	0.211

Results of the HITS algorithm

- When we have reached convergence, we have:

$$h = c_h A a \text{ and } a = c_a A^T h$$

where c_h and c_a are scalars.

i.e. $h = \lambda_h A A^T h$ and $a = \lambda_a A^T A a$

where h is the eigenvector of $A A^T$, and λ_h is the eigenvalue.

Likewise, a is the eigenvector of $A^T A$, and λ_a eigenvalue.

Results of the HITS algorithm

- Check:

$$AA^T h = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} 1 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{bmatrix} \begin{bmatrix} 0.793 \\ 0.566 \\ 0.226 \end{bmatrix} = \begin{bmatrix} 3.732 \\ 2.729 \\ 0.998 \end{bmatrix} = 4.73 \begin{bmatrix} 0.793 \\ 0.566 \\ 0.226 \end{bmatrix}$$

i.e. h is an eigenvector of AA^T , with corresponding eigenvalue 4.73.

Using the HITS algorithm

- Unlike PageRank, the hubs and authorities scores are **computed at query time**
- Given a query q :
 - the **root set** are all documents that include at least one query word
 - the **base set** are all documents linking to, or linked from, a root document
- We then use the base set for computing the hubs and authorities scores

Tolerant retrieval

- Spelling correction
 - “*see you on the wki*” → “*see you on the wiki*”
- Wildcard queries
 - “*colo*rful*” → “*colorful*”, “*colourful*”
- In both cases, the search engine needs to
 - construct the intended query (or queries)
 - compute the results for those queries (intersection, phrase, ranked retrieval)
 - list the results

Wildcard queries: one word

- **care***: find all docs containing any word beginning “care”.
- ***less**: find words ending in “less”
- **colo*r**: find all words beginning “colo” and ending in “r”
- general case: any numbers of ‘*’ placed anywhere in the word (we will **not** consider this case)
- special case: ‘*’ matches all words (you don’t need to consider this case)

Wildcard queries: several words

- **b* color*r:** find all docs containing any word beginning “b”, **and** any word beginning with “color” ending in “r”

Wildcard queries

- How do we find all words matching **care*** ?
- **Idea:** Go through all words in the vocabulary, and check which words match the regular expression
^care.*
 - e.g. using Java's regex library
- Would this work?

K-gram index


- For both wildcard queries and spelling correction we must **quickly** find words that
 - the user intended (for wildcard queries), or
 - the user probably intended (for spelling correction)
- A ***k*-gram index** is an index from *k*-grams (parts of words) to words.
 - **bigram** index when *k*=2
 - **trigram** index when *k*=3
 - etc.

K-grams

- The bigrams of **december**:
 - First add start and end symbol: **^december\$**
 - Bigrams are all two-letter sequences:
^d, de, ec, ce, em, mb, be, er, r\$
- The trigrams of **december**:
 - **^de, dec, ece, cem, emb, mbe, ber, er\$**
- A word of length n has $n+3-k$ k -grams


K-gram index

- For k-gram indexes, we can reuse a lot of ideas from our usual inverted indexes:
 - keys in a hashtable
 - values as arraylists




K-gram index and wildcards

- Suppose we want to find matches for **mon***
- How would you search?




K-gram index and wildcards

- Suppose we want to find matches for **mon***
- How would you search?
 - do an **intersection** search for **^m mo on** in k-gram index
 - post-process the results using the regex library
 - do a **union** search on the resulting words in the ordinary index



K-gram index and wildcards

- Suppose we want to find matches for
mon* tra*
- Which documents in this Venn diagram are you looking for?



Spelling correction

- Two principal uses
 - Correcting document(s) being indexed
 - Correcting user queries to retrieve “right” answers
 - Usually documents are left intact, but queries spell-checked
- Two main flavors:
 - Isolated word
 - Will not catch typos resulting in correctly spelled words, e.g., *from* → *form*
 - Context-sensitive, e.g. ***I flew form Heathrow to Narita.***

Spelling correction

- Fundamental premise – there is a lexicon from which the correct spellings come
- Two basic choices for this
 - **(Grammatical approach)** A standard lexicon such as
 - Webster’s English Dictionary
 - An “industry-specific” lexicon – hand-maintained
 - **(Data-driven approach)** The lexicon of the indexed corpus
 - E.g., all words on the web
 - All names, acronyms etc.
 - (Including the mis-spellings)

Spelling correction of a single word

- What we will do in assignment 3:
 - Data-driven approach (no lexicon)
 - Assumption: A word is **not** misspellt if it appears in at least 1 document.
 - If a word has 0 occurrences it **might** be misspellt, and the search engine should suggest corrections

Spelling correction of a single word

- Methods:
 - **Edit distance**
 - **Weighted edit distance**
 - ***n*-gram overlap**
- These can (should?) be combined

Levenshtein (edit) distance

- What is $\text{dist}(\textit{intention}, \textit{execution})$?

i n t e n t i o n	← delete <i>i</i>
n t e n t i o n	← substitute <i>n</i> by <i>e</i>
e t e n t i o n	← substitute <i>t</i> by <i>x</i>
e x e n t i o n	← insert <i>u</i>
e x e n u t i o n	← substitute <i>n</i> by <i>c</i>
e x e c u t i o n	

- Cost $1+2+2+1+2 = 8$
- Can be efficiently computed with dynamic programming

Computing Levenshtein distance

e							
k							
o							
r							
b							
#							
	#	a	b	o	v	e	

Computing Levenshtein distance

e	5					
k	4					
o	3					
r	2					
b	1					
#	0	1	2	3	4	5
	#	a	b	o	v	e

Computing Levenshtein distance

e	5					
k	4					
o	3					
r	2					
b	1	2				
#	0	1	2	3	4	5
	#	a	b	o	v	e

Computing Levenshtein distance

e	5					
k	4					
o	3					
r	2					
b	1	2	1			
#	0	1	2	3	4	5
	#	a	b	o	v	e

Computing Levenshtein distance

e	5	6	5	4	5	4
k	4	5	4	3	4	5
o	3	4	3	2	3	4
r	2	3	2	3	4	5
b	1	2	1	2	3	4
#	0	1	2	3	4	5
	#	a	b	o	v	e

Weighted edit distances

- As above, but the weight of an operation depends on the character(s) involved
 - Meant to capture OCR or keyboard errors, e.g. m more likely to be mis-typed as n than as q
 - Therefore, replacing m by n is a smaller edit distance than by q
 - This may be formulated as a probability model
- Requires weight matrix as input
- Modify dynamic programming to handle weights

Using edit distances

- Given query, enumerate all strings within a preset (weighted) edit distance (e.g., 2)
- Intersect this set with list of “correct” words
 1. Show terms you found to user as suggestions, ***or***
 2. Look up all possible corrections in our inverted index and return all docs ... slow, ***or***
 3. Run with a single most likely correction
- In assignment 3, we will opt for alternative 1.

Using edit distances

- Given a query, do we compute its edit distance to every dictionary term?
 - Expensive and slow
- How do we find the candidate dictionary terms?
 - One alternative: k-gram overlap
 - Use the k-gram index again!
 - Can also be used by itself for spelling correction

k -gram overlap

- Enumerate all the k -grams in the query string as well as in the lexicon
 - *november*: ^no, nov, ove, vem, emb, mbe, ber, er\$
 - *december*: ^de, dec, ece, cem, emb, mbe, ber, er\$
 - overlap 4/12 unique trigrams
 - the **Jacquard coefficient** = 4/12 = 0.33
- generally,
$$\frac{|X \cap Y|}{|X \cup Y|}$$
 where X,Y are sets

Spelling correction of a single word

- E.g. **wki**
 - Do a union search in the k-gram index for
^w wk ki i\$
 - Calculate the Jacqard coefficient between **wki** and each of the resulting words
 - If the JC > some threshold for word w , calculate the Levenshtein distance between w and **wki**
 - If Levenstein distance < some other threshold, then w is a potential correction
 - Add w to the list of corrections

Spelling correction, multi-word queries

- E.g. "**S**ee **yuoq** on **the** **wki**"
- Includes two (possibly) misspellt words: **yuoq** and **wki** with 0 postings
- Construct the lists of spelling suggestions for each word
 - Lists for words with > 0 postings will only contain themselves
 - Then merge the lists

Spelling correction, multi-word queries

- See youq on the wki

see	you	on	the	wiki
	your			ki
	youd			wi
	yous			wk
	youn			waki

- Now the lists have to merged to produce final suggestions

Spelling correction, multi-word queries

- Final list of suggestions (for instance):
 - See you on the wiki
 - See you on the ki
 - See you on the wi
 - See you on the wk
 - See you on the waki

General issues in spell correction

- We enumerate several possible alternatives to misspelled queries – which ones should we present to the user?
- Use heuristics:
 - The alternative matching most documents (expensive)
 - **The alternative likely to match most documents** (using heuristics, cheaper)
 - Query log analysis – what have others been searching for?
What has this user been searching for?

Assignment 3

- 3.3 K-gram index
- 3.4 Wildcard queries
- 3.5 Spell checking, isolated queries
- 3.6 Spell checking, multi-word queries